Sora 2 on Atlas Cloud: Complete API Guide

Atlas Cloud Team

2026-02-20

Sora 2 on Atlas Cloud: Complete API Guide with Code Examples

OpenAI Sora 2 represents a significant step forward in AI video generation, particularly in the domain of physics simulation. Objects fall, bounce, shatter, and interact with their environments in ways that feel genuinely plausible – a capability that no competing model has fully matched. This Sora 2 tutorial covers everything developers need to integrate Sora 2 into their workflows via the Sora 2 Atlas Cloud API, including pricing, Python code examples, prompt strategies, and a direct comparison against Seedance 2.0, Kling 3.0, and Veo 3.1.

See Sora 2 in action:

- 5 Top AI Models. 3 Prompts. One Clear Winner for Audio/Video Sync – Includes Sora 2 comparison
- OpenAI announces Sora 2 – "the GPT-3.5 moment for video"

The Sora 2 API is available through Atlas Cloud at $0.15 per second of generated video, with $1 in free credit upon signup. Users gain access to Sora 2 alongside 300+ other models through a single API key, eliminating the need to manage separate accounts with OpenAI, Kuaishou, ByteDance, and Google.

Sora 2 at a Glance

Spec	Detail
Developer	OpenAI
Model ID	`openai/sora-2/text-to-video-pro-developer`
Max Resolution	1080p/30fps
Max Duration	12 seconds
Frame Rate	30fps
Native Audio	Yes
Reference Input	1 image
Core Strength	Physics simulation, realistic object interactions
Atlas Cloud Price	$0.15/sec

Key Features of Sora 2

Physics Simulation

This is where Sora 2 genuinely separates itself from the competition. Gravity, momentum, fluid dynamics, and collision behavior are rendered with a level of accuracy that other models cannot consistently replicate. In testing, a prompt describing a bowling ball striking pins produced results with correct force distribution, pin scatter patterns, and follow-through motion. Kling 3.0 and Seedance 2.0 tend to approximate these interactions; Sora 2 simulates them.

Realistic Object Interactions

Beyond simple physics, Sora 2 handles complex multi-object scenes with impressive coherence. Pouring liquid into a glass, stacking blocks that wobble and topple, or a cat knocking objects off a table – these scenarios require the model to understand not just motion, but cause and effect. Sora 2 handles these chains of interaction more reliably than any other model currently available through public APIs.

Material and Surface Rendering

Glass, metal, water, fabric, and wood all exhibit distinct physical properties in Sora 2 outputs. Light refracts through transparent materials. Reflective surfaces catch and distort their surroundings. Cloth drapes and folds with appropriate weight. This attention to material behavior contributes to the overall sense of physical plausibility that defines Sora 2's output.

Native Audio Generation

Sora 2 generates synchronized audio alongside video. Footsteps match walking cadence, impacts produce appropriate sounds, and ambient noise corresponds to the scene. While the audio is not production-grade, it provides a usable starting point that eliminates the need for a separate audio generation step in many workflows.

Temporal Coherence

Over its 12-second maximum duration, Sora 2 maintains strong consistency in lighting, character appearance, and environmental detail. Objects do not spontaneously change color or shape between frames. Shadows track light sources correctly across the full clip. This consistency is particularly important for scenes involving slow, deliberate motion where any discontinuity would be immediately noticeable.

Sora 2 Pricing

OpenAI Sora 2 Official Access

OpenAI Sora 2 is available through ChatGPT with a Plus subscription ($20/month) or Pro subscription ($200/month), but with significant limitations. Plus users receive a restricted number of generations per month, and queue times during peak hours can be substantial. Direct API access through OpenAI requires developer account approval and Sora 2 pricing through this channel is at a premium.

Atlas Cloud API Pricing (Recommended)

Atlas Cloud provides the most straightforward path to the Sora 2 API for developers:

Detail	Value
Model	`openai/sora-2/text-to-video-pro-developer`
Price	$0.15/sec
10-second clip	$1.50
12-second clip (max)	$1.80
Free signup credit	$1.00
Queue	No wait times

A $1 free credit upon registration translates to approximately 6 seconds of Sora 2 video – enough to evaluate quality and physics accuracy before committing to paid usage. This Sora 2 pricing structure means costs scale linearly and predictably for teams generating at volume.

Access Sora 2 API on Atlas Cloud – $1 Free Credit

How to Access Sora 2 API

Option 1: OpenAI Direct

Users can access Sora 2 through OpenAI's official channels, either via ChatGPT subscriptions or the developer API. The developer API requires an approved account and carries higher per-generation costs. Documentation is available at platform.openai.com, though the onboarding process involves waitlist approval for video generation endpoints.

Option 2: Atlas Cloud (Recommended)

For most developers, the Sora 2 Atlas Cloud integration provides the fastest path to production. A single API key unlocks OpenAI Sora 2 alongside 300+ other models, including Seedance 2.0, Kling 3.0, and Veo 3.1. No waitlist. No separate accounts. Unified billing.

Step 1: Create an account at atlascloud.ai and retrieve your API key from the dashboard. The $1 free credit is applied automatically.

Step 2: Use the following Python code to generate video with Sora 2:

```python
import requests
import time

API_KEY = "your-atlas-cloud-api-key"
BASE_URL = "https://api.atlascloud.ai/api/v1"

# Generate video with Sora 2
response = requests.post(
    f"{BASE_URL}/model/generateVideo",
    headers={
        "Authorization": f"Bearer {API_KEY}",
        "Content-Type": "application/json"
```

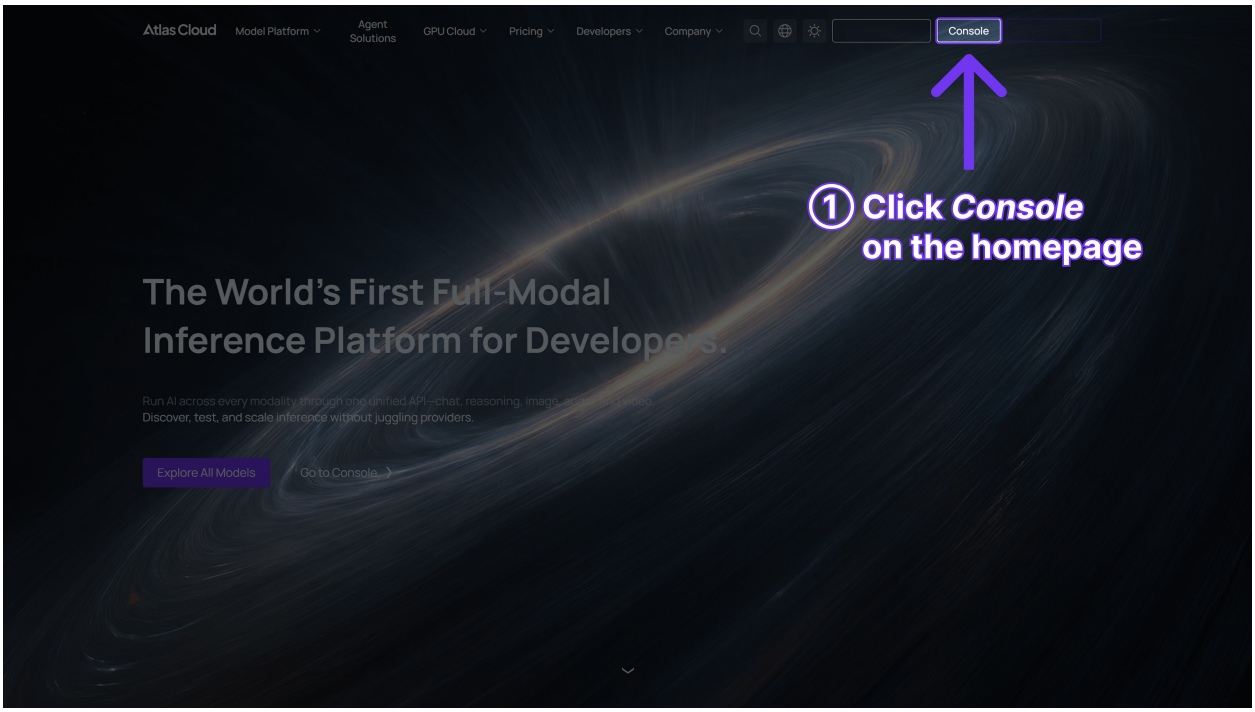

Figure 1: How to create an API key on Atlas Cloud

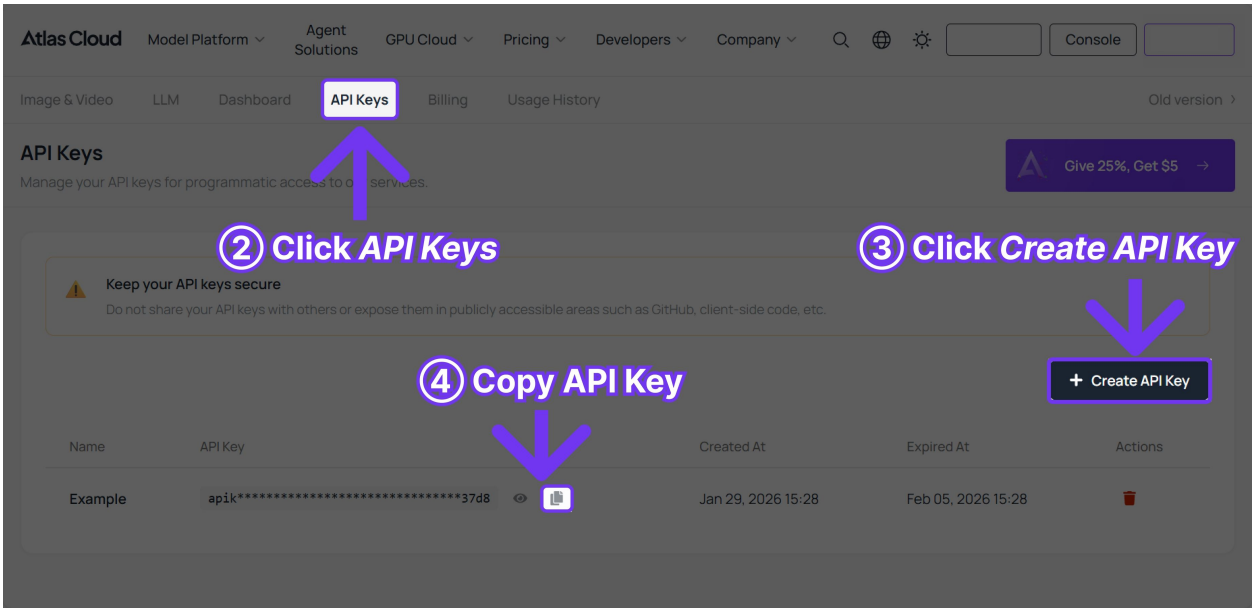

Figure 2: API key management on Atlas Cloud console

```
    },
    json={
        "model": "openai/sora-2/text-to-video-pro-developer",
        "prompt": "A glass sphere rolling down a wooden staircase, each bounce creating ripples
        "duration": 10,
        "resolution": "1080p"
    }
)

result = response.json()

# Poll for results
while True:
    status = requests.get(
        f"{BASE_URL}/model/prediction/{result['request_id']}/get",
        headers={"Authorization": f"Bearer {API_KEY}"}
    ).json()
    if status["status"] == "completed":
        print(f"Video: {status['output']['video_url']}")
        break
    time.sleep(5)
```

Step 3: The API returns a `request_id` immediately. Poll the prediction endpoint until the status changes to `completed`, then retrieve the video URL from the response. Generation times for Sora 2 typically range from 30 seconds to 3 minutes depending on duration and complexity.

Start Using Sora 2 on Atlas Cloud

Sora 2 Prompt Tips

After extensive testing with the Sora 2 API, this Sora 2 tutorial section covers several prompting strategies that consistently produce superior results. Sora 2 responds particularly well to prompts that describe physical interactions and material behavior.

1. Describe physics explicitly. Sora 2's greatest strength is physical simulation, so lean into it. Rather than "a ball bouncing," specify "a rubber ball dropped from shoulder height onto a hardwood floor, bouncing three times with decreasing height, realistic elasticity and shadow movement." The model rewards specificity in physical descriptions.

2. Specify materials and surfaces. Sora 2 renders different materials with distinct properties. Include material descriptors: "frosted glass," "brushed aluminum," "wet cobblestone," "silk curtain." The model differentiates between these textures and applies appropriate light behavior to each.

3. Use cinematic language for camera work. Sora 2 interprets standard cinematography terminology with good accuracy. "Slow dolly forward," "rack focus from foreground to background," and "low-angle tracking shot" all produce the expected camera behavior. Avoid vague directions like "cool camera movement."

4. Design prompts around 10-12 seconds. With a maximum duration of 12 seconds, prompts

should describe a single coherent action or sequence. One clear subject, one primary motion, one visual payoff. Avoid cramming multiple scenes or transitions into a single generation.

5. Leverage cause-and-effect chains. Sora 2 handles sequential physical events well. "A domino chain reaction across a desk, knocking over a cup of pencils" or "a match striking, igniting, and lighting a candle wick" – these multi-step physical sequences play to the model's strengths.

Example prompts that performed well in testing:

Physics showcase:

```
A ceramic mug tipping off the edge of a kitchen counter in slow motion,
shattering on impact with a tile floor, fragments and liquid scattering
realistically, soft morning light from a nearby window, cinematic depth of field
```

Product demonstration:

```
A smartphone placed on a reflective dark surface, water droplets falling
onto the screen and beading naturally, demonstrating water resistance,
studio lighting with subtle blue accent lights, 4K commercial style
```

Nature simulation:

```
A single drop of water falling into a still pond, creating concentric
ripples that expand outward, a fallen autumn leaf floating on the surface
gently disturbed by the waves, golden hour lighting, macro lens perspective
```

Sora 2 vs Competitors

Feature	Sora 2	Seedance 2.0	Kling 3.0	Veo 3.1
Max Resolution	1080p/30fps	2K/24fps	4K/60fps	1080p/24fps
Max Duration	12s	15s	10s	8s
Reference Input	1 image	12 files	1-2 images	1-2 images
Native Audio	Yes	Yes	Yes (5 languages)	Yes
API Cost (Atlas Cloud)	$0.15/sec	$0.022/sec	$0.126/sec	$0.03/sec
Best Strength	Physics realism	Multimodal control	Resolution + value	Cinematic polish
Content Filter	Strict	Strict	Very strict	Moderate

Where Sora 2 Wins

OpenAI Sora 2 leads the field in physics simulation accuracy and realistic object interaction. For content that involves physical cause and effect – product demonstrations, educational visualizations, physics-driven creative work – no other model consistently matches its output quality. The 12-second maximum duration also gives it an edge over Kling 3.0 (10s) and Veo 3.1 (8s), providing more room for narrative development within a single clip.

Where Sora 2 Falls Short

When evaluating Sora 2 pricing, at $0.15/sec Sora 2 is the most expensive option in this comparison. Seedance 2.0 at $0.022/sec costs roughly 85% less per second. Resolution is capped at 1080p/30fps, which lags behind Kling 3.0's 4K/60fps capability. Reference input is limited to a single image, whereas Seedance 2.0 supports up to 12 files (nine images, three videos, three audio). For teams on tight budgets or projects requiring extensive reference-based control, these limitations are significant.

The Practical Approach

As this Sora 2 tutorial demonstrates, most production teams will benefit from having access to multiple models rather than committing exclusively to Sora 2. Physics-heavy scenes go to Sora 2. Complex multi-reference projects go to Seedance 2.0. High-resolution final renders go to Kling 3.0. Cinematic polish goes to Veo 3.1. Atlas Cloud makes this multi-model workflow practical with a single API key and unified billing.

Who Should Use Sora 2?

Choose Sora 2 if:

- Physics accuracy is the primary requirement. Gravity, fluid dynamics, collisions, and material interactions are Sora 2's defining strength.
- The project involves product demonstrations where objects need to behave realistically – drops, splashes, rotations, mechanical movements.
- Educational or scientific visualization work demands plausible physical behavior that withstands scrutiny.
- The 12-second duration ceiling provides sufficient time for the intended content. Sora 2 offers more time per clip than Kling 3.0 and Veo 3.1.
- Budget is less of a concern than output quality for physics-driven content.

Choose Seedance 2.0 instead if:

- The project requires extensive reference material – multiple images, videos, and audio files as input.
- Budget efficiency is paramount. At $0.022/sec on Atlas Cloud, Seedance 2.0 is roughly 7x cheaper than Sora 2.
- Longer clips (up to 15 seconds) are needed.
- Multimodal input control is more important than physics accuracy.

Choose Kling 3.0 instead if:

- 4K/60fps output is required. Sora 2 caps at 1080p/30fps.
- Free-tier access matters. Kling 3.0 offers 66 daily credits; Sora 2 requires a paid subscription.
- E-commerce content needs crisp, readable text rendering in generated video.

Choose Veo 3.1 instead if:

- Cinematic visual quality and color grading are the top priorities.
- Cost-effective scaling is needed. At $0.03/sec, Veo 3.1 is 5x cheaper than Sora 2.
- Shorter clips (up to 8 seconds) are acceptable for the use case.

Frequently Asked Questions

How do I access the Sora 2 API?

The most straightforward method is through Atlas Cloud. Create an account, retrieve an API key, and use the model ID `openai/sora-2/text-to-video-pro-developer` in your requests. The $1 free credit is applied automatically at signup. Alternatively, Sora 2 is accessible through OpenAI's official API, though this requires separate developer account approval.

How much does Sora 2 cost per video?

Sora 2 pricing on Atlas Cloud is $0.15 per second of generated video. A 10-second clip costs $1.50. A maximum-length 12-second clip costs $1.80. Through OpenAI directly, Sora 2 pricing varies by subscription tier, but generally runs higher than the Sora 2 Atlas Cloud rate for equivalent output.

What is the maximum video length Sora 2 can generate?

Sora 2 supports a maximum duration of 12 seconds at 1080p/30fps. This positions it between Seedance 2.0 (15 seconds maximum) and Kling 3.0 (10 seconds maximum). For longer content, users typically generate multiple clips and stitch them together in post-production.

Can Sora 2 do image-to-video generation?

Yes. Sora 2 accepts a single reference image as input, which the model uses as a starting frame or style reference for the generated video. However, it is limited to one image. Seedance 2.0 supports up to 12 reference files for more complex reference-based generation.

Is Sora 2 better than Seedance 2.0 or Kling 3.0?

Each model excels in different areas. Sora 2 leads in physics simulation and realistic object interactions. Seedance 2.0 offers superior multimodal input control and the lowest API pricing at $0.022/sec. Kling 3.0 provides the highest resolution at 4K/60fps and the most generous free tier. The best approach for most teams is to use all three through a unified platform like Atlas Cloud, selecting the right model for each specific task.

Does Sora 2 generate audio with the video?

Yes. Sora 2 includes native audio generation that is synchronized with the visual content. The audio quality is serviceable for drafts and social media content, though professional productions may benefit from replacing or enhancing the generated audio in post-production.

Verdict

Sora 2 occupies a distinct position in the AI video generation landscape. It is not the cheapest option, nor does it offer the highest resolution or the most flexible input system. What it delivers is the most physically accurate video generation available through any public API. For developers and creators whose work depends on believable physics – product demonstrations, educational content, scientific visualization, or creative work grounded in real-world physical behavior – Sora 2 is the strongest choice available.

The practical recommendation from this Sora 2 tutorial remains consistent: access Sora 2 through Atlas Cloud alongside Seedance 2.0, Kling 3.0, Veo 3.1, and 300+ other models. One API key. One bill. The flexibility to choose the right model for each project. The $1 free credit provides enough to evaluate Sora 2's physics simulation firsthand before scaling up.

Get $1 Free Credit on Atlas Cloud – Try Sora 2 and 300+ Models

---

Related Articles

- Seedance 2.0 Pricing Breakdown
- Kling 3.0 Review: Features, Pricing & How to Access
- Veo 3.1 on Atlas Cloud: Google's 4K AI Video Guide
- How to Use Seedance 2.0 for Video Generation
- 15 Best Seedance 2.0 Prompts for Viral Videos